

# EECS 442 Computer Vision: Homework 5

## Instructions

- This homework is due **due at 5:00:00 p.m. on April 17, 2023.**
- The submission includes two parts:

1. **To Canvas:** submit a zip file of all of your code.

**We have indicated questions where you have to do something in code in red.**

**We have indicated questions where we will definitely use an autograder in purple.**

Please be especially careful on the autograded assignments to follow the instructions. Don't swap the order of arguments and do not return extra values.

If we're talking about autograding a filename, we will be pulling out these files with a script. Please be careful about the name.

Your zip file should contain a single directory which has the same name as your unickname. If I (David, unickname fouhey) were submitting my code, the zip file should contain a single folder fouhey/ containing all required files.

**What should I submit? At the end of the homework, there is a canvas submission checklist provided.** We provide a script that validates the submission format [here](#). If we don't ask you for it, you don't need to submit it; while you should clean up the directory, don't panic about having an extra file or two.

2. **To Gradescope:** submit a pdf file as your write-up, including your answers to all the questions and key choices you made.

**We have indicated questions where you have to do something in the report in blue.**

You might like to combine several files to make a submission. Here is an example online link for combining multiple PDF files: <https://combinepdf.com/>.

The write-up must be an electronic version. **No handwriting, including plotting questions.** L<sup>A</sup>T<sub>E</sub>X is recommended but not mandatory.

## Python Environment

We are using Python 3.7 for this course. You can find references for the Python standard library here: <https://docs.python.org/3.7/library/index.html>. To make your life easier, we **recommend** you to install the latest Anaconda for Python 3.7 (<https://www.anaconda.com/download/>). This is a Python package manager that includes most of the modules you need for this course.

We will make use of the following packages extensively in this course:

- Numpy (<https://docs.scipy.org/doc/numpy-dev/user/quickstart.html>).
- Matplotlib ([http://matplotlib.org/users/pyplot\\_tutorial.html](http://matplotlib.org/users/pyplot_tutorial.html)).
- OpenCV (<https://opencv.org/>).

# 1 Camera Calibration

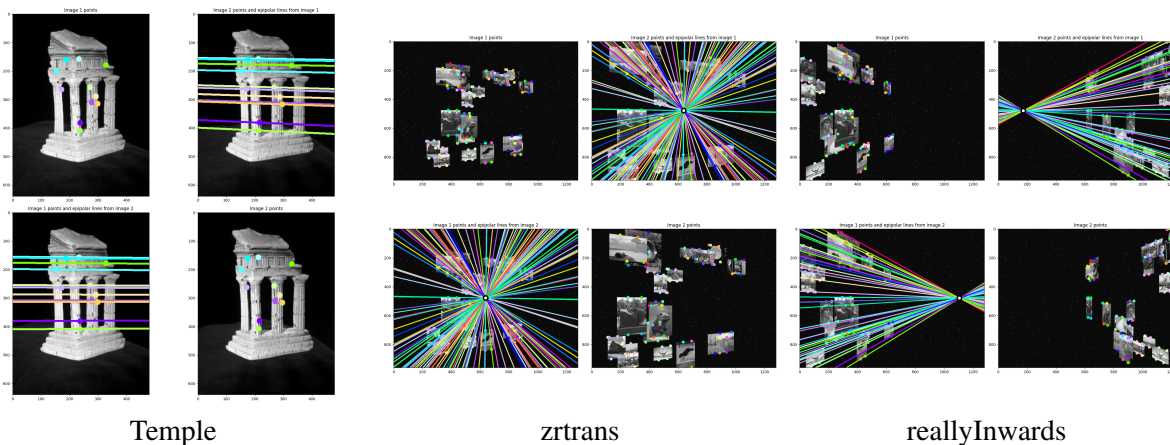


Figure 1: Epipolar lines for some of the datasets.

## Task 1: Estimating $M$ [35 points]

We will give you a set of 3D points  $\{\mathbf{X}_i\}_i$  and corresponding 2D points  $\{\mathbf{p}_i\}_i$ . The goal is to compute the projection matrix  $M$  that maps from world 3D coordinates to 2D image coordinates. Recall that

$$\mathbf{p} \equiv M\mathbf{X}, \quad (1)$$

and (see foreword) by deriving an optimization problem. The script `task1.py` shows you how to load the data. The data we want you to use is in `task1/`, but we show you how to use data from Task 2 and 3 as well. **Credit:** The data from task 1 and an early version of the problem comes from James Hays’s Georgia Tech CS 6476.

- (a) (15 points) **Fill in `find_projection` in `task1.py`.**
- (b) (5 points) **Report  $M$**  for the data in `task1/`.
- (c) (10 points) **Fill in `compute_distance` in `task1.py`.** In this question, you need to compute the average distance in the image plane (i.e., pixel locations) between the homogeneous points  $M\mathbf{X}_i$  and 2D image coordinates  $\mathbf{p}_i$ , or

$$\frac{1}{N} \sum_i^N \|\text{proj}(M\mathbf{X}_i) - \mathbf{p}_i\|_2. \quad (2)$$

where  $\text{proj}([x, y, w]) = [x/w, y/w]$ . The distance quantifies how well the projection maps the points  $\mathbf{X}_i$  to  $\mathbf{p}_i$ . You should use `find_projection` from part a). Note: You should feel good about the distance if it is **less than 0.01** for the given sample data. If you plug in different data, this threshold will of course vary.

- (d) (5 points) **Describe what relationship, if any, there is between Equation 2 as above and Equation 6 in the HW5 Notes** Note that the points we’ve given you are well-described by a linear projection – there’s no noise in the measurements – but in practice, there will be an error that has to minimize. Both equations represent objectives that could be used. If they are the same, show it; if they are not the same, report which one makes more sense to minimize. Things to consider include whether the equations directly represent anything meaningful.

## 2 Estimation of the Fundamental Matrix and Reconstruction

**Data:** we give you a series of datasets that are nicely bundled in the folder `task23/`. Each dataset contains two images `img1.png` and `img2.png` and a numpy file `data.npz` containing a whole bunch of variables. The script `task23.py` shows how to load the data.

**Credit:** `temple` comes from Middlebury's Multiview Stereo dataset. The images shown in the synthetic images are described in HW1's credits.

### Task 2: Estimating $F$ [35 points]

- (a) (15 points) **Fill in `find_fundamental_matrix`** in `task23.py`. You should implement the eight-point algorithm. Remember to normalize the data and to reduce the rank of  $F$ . For normalization, you can scale the image size and center the data at 0.
- (b) (10 points) **Fill in `compute_epipoles`**. This should return the homogeneous coordinates of the epipoles – remember they can be infinitely far away!
- (c) (5 points) **Show epipolar lines for `temple`, `reallyInwards`, and another dataset of your choice**.
- (d) (5 points) **Report the epipoles for `reallyInwards` and `xtrans`**.

### Task 3: Triangulating X [30 points]




Figure 2: Visualizations of **reallyInwards** reconstructions

The next step is extracting 3D points from 2D points and camera matrices, which is called triangulation. Let  $\mathbf{X}$  be a point in 3D.

$$\mathbf{p} = \mathbf{M}_1 \mathbf{X} \quad \mathbf{p}' = \mathbf{M}_2 \mathbf{X} \quad (3)$$

Triangulation solves for  $\mathbf{X}$  given  $\mathbf{p}, \mathbf{p}', \mathbf{M}_1, \mathbf{M}_2$ . We'll use OpenCV's algorithms to do this.

- (a) (5 points) **Compute the Essential Matrix  $\mathbf{E}$  for the Fundamental Matrix  $\mathbf{F}$ .** You should do this for the dataset `reallyInwards`. Recall that

$$\mathbf{F} = \mathbf{K}'^{-T} \mathbf{E} \mathbf{K}^{-1} \quad (4)$$

and that  $\mathbf{K}, \mathbf{K}'$  are always invertible (for reasonable cameras), so you can compute  $\mathbf{E}$  straightforwardly.

- (b) (15 points) **Fill in `findTriangulation` in `task23.py`.**

The first camera's projection matrix is  $\mathbf{K}[\mathbf{I}, \mathbf{0}]$ . The second camera's projection matrix can be obtained by decomposing  $\mathbf{E}$  into a rotation and translation via `cv2.decomposeEssentialMat`. (Note:  $\mathbf{E}$  can be obtained using the formula from part a) This function returns two matrices  $\mathbf{R}_1$  and  $\mathbf{R}_2$  and a translation  $\mathbf{t}$ . The four possible camera matrices for  $\mathbf{M}_2$  are:

$$\mathbf{M}_2^1 = \mathbf{K}'[\mathbf{R}_1, \mathbf{t}], \quad \mathbf{M}_2^2 = \mathbf{K}'[\mathbf{R}_1, -\mathbf{t}], \quad \mathbf{M}_2^3 = \mathbf{K}'[\mathbf{R}_2, \mathbf{t}], \quad \mathbf{M}_2^4 = \mathbf{K}'[\mathbf{R}_2, -\mathbf{t}] \quad (5)$$

You can identify which projection is correct by picking the one for which the most 3D points are in front of both cameras. You'll have to run `cv2.triangulatePoints` in order to do this. Once you have the 3D point  $\mathbf{X}$  and a potential camera matrix  $\mathbf{M}_2$ , you can compute  $\mathbf{M}_2 \mathbf{X}$  and then check the last entry  $w$  of the resulting homogeneous coordinate  $[u, v, w] = \mathbf{M}_2 \mathbf{X}$ . The reason why this works is that the extrinsics put the 3D point in the camera's frame, where  $z < 0$  is behind the camera, and the last row of  $\mathbf{K}$  is  $[0, 0, 1]$  so this does not change things.

Then, triangulate the 2D points using `cv2.triangulatePoints` using the  $\mathbf{M}_2$  that makes the most points positive.

*Note:* you can approach this in two ways: (1) try all the  $M_2$  and figure out which one puts the most points in front of the camera; then do the triangulation to get the final set of points; or (2) try all the  $M_2$  options and keep track of both the “best” set of points and the  $M_2$ , and then return “the best” set of points. Either is fine.

- (c) (10 points) **Put a visualization of the point cloud for `reallyInwards` in your report.** You can use `visualize_pcd` in `utils.py` or implement your own.

## References and Credits

- Temple dataset used in Tasks 2 and 3: <http://vision.middlebury.edu/mview/data/>.
- Part of the homework are taken from Georgia Tech CS 6476 by James Hays and CMU 16-385. Please feel free to similarly re-use our problems while similarly crediting us.